

Comparing ROC Packages

June 19, 2020

1 Packages

```
[16]: library(mdsr)
library(rpart)
library(nnet)
library(e1071)
library(randomForest)
library(ROCR)
library(plotROC)
library(ROCit)
```

2 Init

```
[2]: census <- read.csv(
  "http://archive.ics.uci.edu/ml/machine-learning-databases/adult/adult.data",
  header = FALSE)
names(census) <- c("age", "workclass", "fnlwgt", "education",
  "education.num", "marital.status", "occupation",
  "relationship", "race", "sex", "capital.gain",
  "capital.loss", "hours.per.week", "native.country",
  "income")

set.seed(364)
n <- nrow(census)
test_idx <- sample.int(n, size = round(0.2 * n))
train <- census[-test_idx,]
test <- census[test_idx,]

form <- as.formula("income ~ age + workclass + education + marital.status +
  occupation + relationship + race + sex + capital.gain + capital.loss +
  hours.per.week")

cut_seq <- seq(0, 1, 0.1)
cut_seq_2 <- seq(0, 1, 0.25)
```

3 Models

3.1 Decision Tree

```
[3]: mod_tree <- rpart(form, data = train)
income_tree_probs <- mod_tree %>% predict(newdata = test, type = "prob")
df_tree <- data.frame(predictions = income_tree_probs[, ' >50K'],
                      labels = ifelse(test$income == ' >50K', 1, 0))
```

3.2 Naive Bayes

```
[4]: mod_nb <- naiveBayes(form, data = train)
income_nb_probs <- mod_nb %>% predict(newdata = test, type = "raw")
df_nb <- data.frame(predictions = income_nb_probs[, ' >50K'],
                   labels = ifelse(test$income == ' >50K', 1, 0))
```

3.3 Neural Network

```
[5]: mod_nn <- nnet(form, data = train, size = 5)
income_nn_probs <- mod_nn %>% predict(newdata = test, type = "raw")
df_nn <- data.frame(predictions = income_nn_probs,
                   labels = ifelse(test$income == ' >50K', 1, 0))
```

```
# weights: 296
initial value 17863.316757
iter 10 value 13160.102775
iter 20 value 11490.083360
iter 30 value 10492.265833
iter 40 value 9190.590748
iter 50 value 8868.952933
iter 60 value 8669.234503
iter 70 value 8582.637394
iter 80 value 8516.139198
iter 90 value 8442.790655
iter 100 value 8408.961720
final value 8408.961720
stopped after 100 iterations
```

3.4 Random Forest

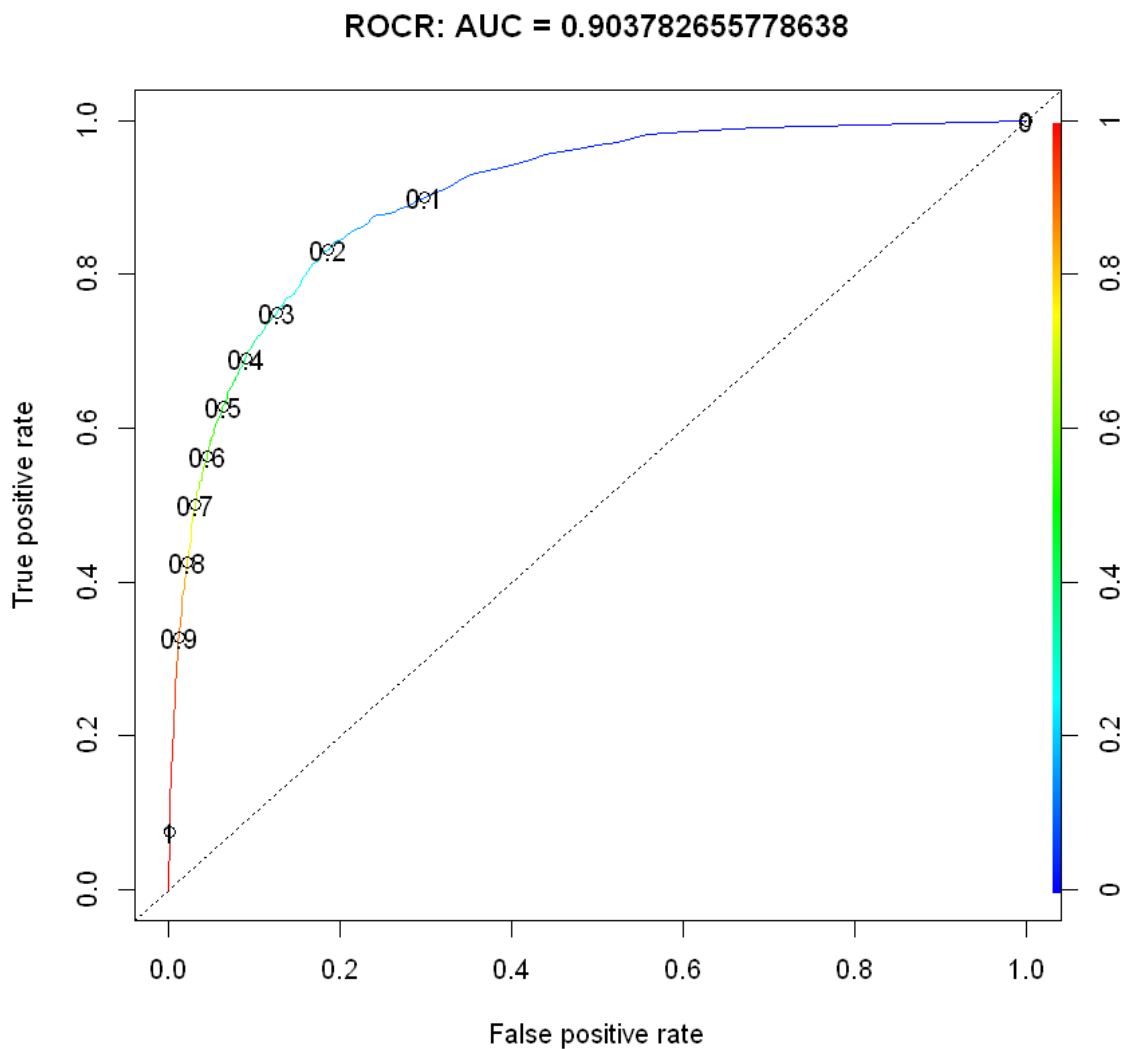
```
[6]: mod_forest <- randomForest(form, data = train, ntree = 201, mtry = 3)
income_forest_probs <- mod_forest %>% predict(newdata = test, type = "prob")
df_forest <- data.frame(predictions = income_forest_probs[, ' >50K'],
                       labels = ifelse(test$income == ' >50K', 1, 0))
```

4 Comparing Packages

4.1 ROCR

ROCR allows you to colorize a ROC curve according to threshold values. On top of that, you can also label a certain set of threshold values on the curve, giving you two different methods of identifying cutoffs.

```
[7]: rocr_pred_forest <- prediction(df_forest$predictions, df_forest$labels)
rocr_perf_forest <- performance(rocr_pred_forest, 'tpr', 'fpr')
rocr_perf_forest %>% plot(colorize = TRUE,
                          print.cutoffs.at = cut_seq)
rocr_auc_forest <- performance(rocr_pred_forest, 'auc')@y.values[[1]]
title(paste("ROCR: AUC =", rocr_auc_forest))
abline(0, 1, lty = 3)
```




Overlaying curves is actually fairly simple with `ROCR`. All you have to do is plot an initial curve, then every subsequent curve you plot must specify `add = TRUE`. This prevents the previous graph from going away. This, combined with `colorize`, provides an interesting way of comparing ROC curves and how their thresholds are distributed differently along them. The drawback of this approach is that there isn't an easy way to label or identify each line, other than remembering the shape of each one.

```
[8]: rocr_pred_tree <- prediction(df_tree$predictions, df_tree$labels)
      rocr_pred_nb <- prediction(df_nb$predictions, df_nb$labels)
      rocr_pred_nn <- prediction(df_nn$predictions, df_nn$labels)

      rocr_perf_tree <- performance(rocr_pred_tree, 'tpr', 'fpr')
      rocr_perf_nb <- performance(rocr_pred_nb, 'tpr', 'fpr')
      rocr_perf_nn <- performance(rocr_pred_nn, 'tpr', 'fpr')

      rocr_perf_forest %>% plot(colorize = TRUE)
      rocr_perf_tree %>% plot(colorize = TRUE, add = TRUE)
      rocr_perf_nb %>% plot(colorize = TRUE, add = TRUE)
      rocr_perf_nn %>% plot(colorize = TRUE, add = TRUE)
      abline(0, 1, lty = 3)
      title("ROCR")
```




We can emulate the same idea as before but with identifiers for each curve. We lose the gradients in exchange for color identifiers and substitute in labels to identify threshold locations. The labels are all black, so it is easy to confuse nearby labels.

```
[9]: rocr_perf_forest %>% plot(print.cutoffs.at = cut_seq_2,
                             col = "red")
rocr_perf_tree %>% plot(print.cutoffs.at = cut_seq_2,
                       col = "orange", add = TRUE)
rocr_perf_nb %>% plot(print.cutoffs.at = cut_seq_2,
                     col = "green", add = TRUE)
rocr_perf_nn %>% plot(print.cutoffs.at = cut_seq_2,
                    col = "blue", add = TRUE)
abline(0, 1, lty = 3)
```

```

legend("right", legend = c("RF", "DT", "NB", "NN"),
      col = c("red", "orange", "green", "blue"),
      lty = 1, cex = 0.8)
title("ROCR")

```




4.2 plotROC

As opposed to `ROCR`, which does not integrate with `ggplot` by default, `plotROC` does. This has its tradeoffs. In order to calculate AUC, you must first have a `ggplot` object with a `geom_roc` layer. This is because there is no `plotROC` object. It also means that there is no easy way of graphing a ROC curve with a color gradient, since the threshold values are calculated behind the scenes and never returned. However, the labeling is a bit nicer with `plotROC`, as they are offset a bit from the curve. Also, the AUC calculations in this package are sometimes different than those in `ROCR` and


ROCit, sometimes starting as early as the fourth or fifth decimal place.

```
[10]: plotroc_plot <- df_forest %>% ggplot(aes(m = predictions, d = labels)) +  
  geom_roc(cutoffs.at = cut_seq) +  
  geom_abline(slope = 1, intercept = 0, lty = 3)  
plotroc_auc <- calc_auc(plotroc_plot)[['AUC']]  
plotroc_ggplot <- plotroc_plot +  
  ggtitle(paste("plotROC: AUC =", plotroc_auc))  
plotroc_ggplot
```



Overlaying multiple curves is fairly simple. All you have to do is combine them all into one data frame and give them an identifier by model, then separate them by color on the plot. The labels are colored according to which model they belong to, making this graph easier to interpret and nicer to look at than the similar ROCit overlay.

```
[11]: df_all <- rbind(mutate(df_forest, mod = "RF"),
                     mutate(df_tree, mod = "DT"),
                     mutate(df_nb, mod = "NB"),
                     mutate(df_nn, mod = "NN"))
df_all %>% ggplot(aes(m = predictions, d = labels, color = mod)) +
  geom_roc(cutoffs.at = cut_seq_2, labelround = 2) +
  geom_abline(slope = 1, intercept = 0, lty = 3)
```

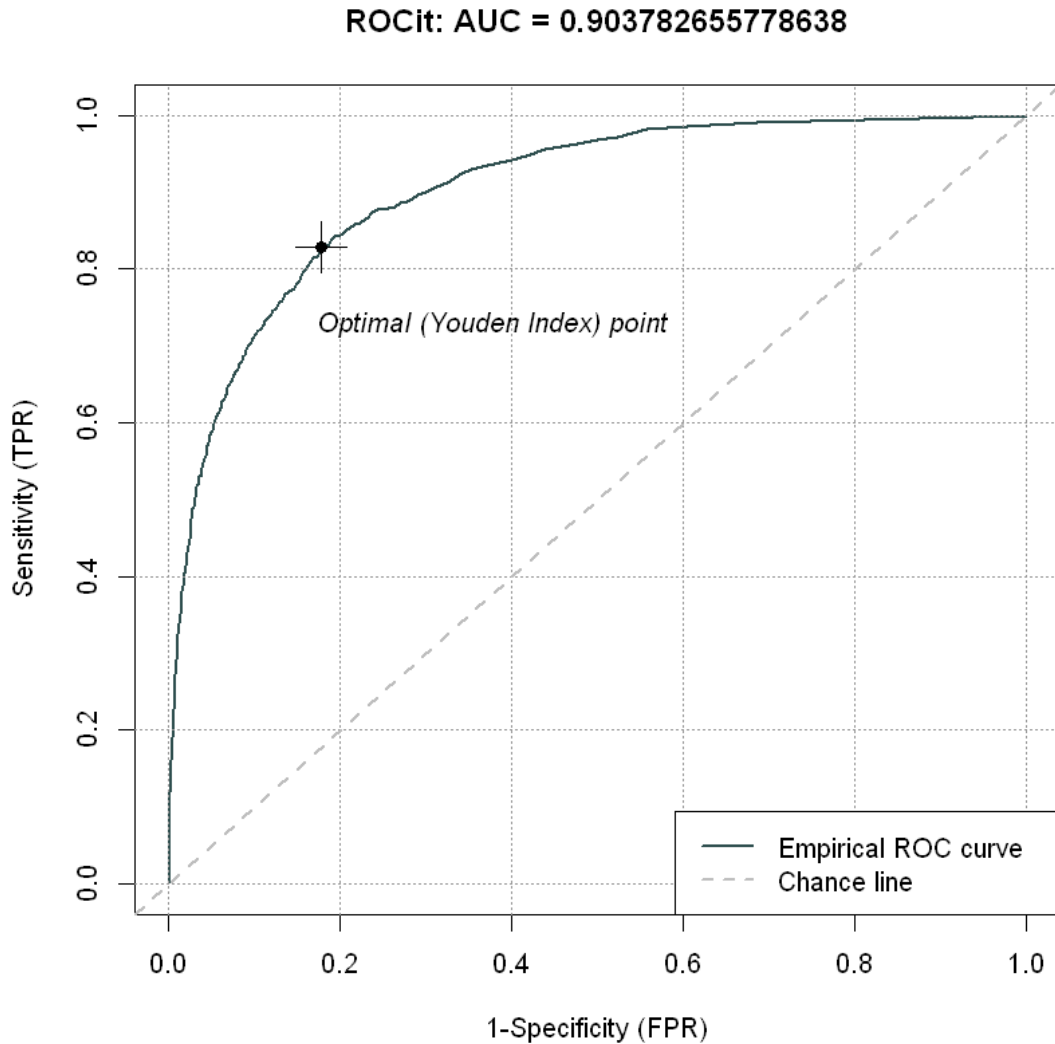


4.3 ROCit

ROCit is interesting because, by default, it plots an “Optimal (Youden Index) point” for each curve, which is the point on the curve that is the farthest from the Chance line vertically. This gives a useful but perhaps loose benchmark for optimality. The `rocit_plot` also silently has all the fpr,

tpr, and cutoff values, but does not display the cutoff values at all.

```
[12]: rocit_obj <- rocit(df_forest$predictions, df_forest$labels)
      rocit_plot <- rocit_obj %>% plot()
      rocit_auc <- rocit_obj[['AUC']]
      title(paste("ROCit: AUC =", rocit_auc))
```



4.4 ROCit and ROCR Combined

If you want the nice look of a ROCit plot with the labelling provided by ROCR, you can combine the two in the same way we combined the multiple ROCR plots from before, by specifying `add = TRUE`.

```
[15]: rocit_obj %>% plot(YIndex = FALSE)
      rocr_perf_forest %>% plot(print.cutoffs.at = cut_seq, add = TRUE)
      title(paste("ROCit + ROCR: AUC =", rocit_auc))
```

